

GALLERY DUTCH ART

AI Agents For Everyone

The practical, hands-on workbook — build and use AI agents yourself

gallerydutchart.nl

Don't just read about AI agents. Build, use and master them — one hands-on exercise at a time.

Written in the spirit of the **For Dummies** series: plain language, real examples, and — above all — ***things you actually do.*** Keep this open next to your screen and complete the exercises as you go.

Icons used in this workbook

- **Tip** — a shortcut or smart move
- **Remember** — the one thing to take away
- **Warning** — how to avoid getting burned
- **'Try It Yourself'** — a hands-on exercise. Don't skip these — they're the whole point.
- **Template** — copy, paste, fill in.
- **Technical Stuff** — skip it if you like; it won't be on a test.

How to use this workbook: Each chapter has a little teaching, then a hands-on exercise. Do the exercise before moving on. By the end you'll have used an agent, written a real brief, and built your first simple agent. Total hands-on time: about 3–4 hours.

Chapter 1: What an AI agent actually is (and a 5-minute first contact)

A regular AI *talks*. An **AI agent** *acts*: it uses tools (a browser, your calendar, email) to finish a multi-step task, deciding its own next step. Picture a brilliant intern who is finally allowed to use the office, not just chat with you.

The four ingredients: a **goal**, a **brain** (the AI model), **tools** (what it can use), and a **loop** (try !' check !' try again).

Remember: Chatbot = talks. Agent = does. That's the whole revolution.

'Try It Yourself — First contact (5 min)

1. Open any modern assistant (ChatGPT, Claude or Gemini).
2. Paste: *"Act as my assistant. Ask me 3 questions, then make a simple weekend plan for me."*
3. Notice how it takes initiative and works in steps. That "taking initiative across steps" is the seed of agent behaviour.
4. Write one sentence below: *What boring task in my week would I most love to hand to an agent?* Keep that task — you'll automate it later.

My task to automate: _____

Chapter 2: Spot the work an agent should do

Not every task suits an agent. The sweet spot is **boring + repetitive + low-risk**. Start there, win, then expand.

'Try It Yourself — Your automation audit (10 min)

List your recurring tasks for one week. For each, score 1–5:

- **Boring** (5 = very), **Repetitive** (5 = daily/weekly), **Low-risk** (5 = nothing breaks if it's wrong).

Template — Task scorecard

Task	Boring	Repetitive	Low-risk	Total
(e.g. summarise newsletters)	5	5	5	15

The task with the **highest total** is your first agent project. Circle it.

Warning: A high-risk task (money, contracts, public posts) is a *bad* first project even if it's boring. Earn trust on safe tasks first.

Chapter 3: The only vocabulary you need

• **Model / LLM** — the AI brain. **Prompt** — your instructions. **Tool / function calling** — letting the AI use an app. **MCP** — a standard "plug" so agents connect to tools safely (USB-C for AI). **Autonomy** — how much it does without asking. **Guardrails** — the rules that stop it doing something silly.

Tip: Judging any "AI agent" product? Ask just two things: *What tools can it use?* and *What can it do without asking me?*

'Try It Yourself' — Two-question test (5 min)

Pick one AI product you've heard of. Write its answers:

- Tools it can use: _____
- What it does without asking: _____

If you can't find out, that's a red flag about its transparency.

Chapter 4: Use your first agent (a 30-minute lab, no code)

Time for a real one. We'll automate the task you circled in Chapter 2.

'Lab 1' — Your first working agent (30 min)

1. **Pick the task** (your circled one).
2. **Pick a tool:** a consumer assistant that can take actions, or a no-code "agent builder" app.
3. **Write the goal as a brief** (use the template below) — not a riddle.
4. **Connect one tool** it needs (calendar, a folder, web search).
5. **Run it once, fully** — approve every step the first time.
6. **Fix the brief** based on what went wrong, and run again.

Template — The agent brief (copy & fill)

ROLE: You are my [assistant for ____].
GOAL: [What "done" looks like, concretely.]
INPUTS: [Where to get information.]
STEPS: [If known, the rough steps.]
RULES: Always [ask before ____]. Never [____].
OUTPUT: [Format you want, e.g. a 5-bullet summary.]
STOP: If unsure, stop and ask me.

Warning: Don't connect your bank, main email or anything destructive on day one. Use a test account or read-only mode.

Remember: A good brief reads like instructions to a new assistant on their first morning.

Chapter 5: Build a simple agent of your own (a 45-minute lab)

You don't need to code. An agent is just a loop: **plan !' act with a tool !' check !' repeat.**

'Lab 2 — Build a two-step agent (45 min)

1. **Map it on paper:** write every step a human takes to do the task.
2. **Mark the tools** each step needs (search, read a sheet, send a draft).
3. **Open a no-code agent builder** (or an assistant that supports custom actions/MCP).
4. **Build step 1 only. Test it.** Then add step 2. Test again.
5. **Add guardrails:** "always ask before sending," a spending cap, a website allow-list.
6. **Add a fallback:** "if unsure, stop and ask."

Template — Workflow map

```
Step 1: _____ !' tool: _____ !' check: _____  
Step 2: _____ !' tool: _____ !' check: _____  
Guardrail: _____ Fallback: if unsure !' ask me
```

Tip: Start with ONE tool (just web search). Add tools only after the single-tool version works.

Technical Stuff: "Agent frameworks" and "SDKs" just package that plan-act-check loop. You can master the *idea* without ever seeing code.

Chapter 6: Stay in control — a safety checklist you'll actually use

Power needs brakes. Run this checklist before you let any agent off the leash.

'Try It Yourself — The pre-flight checklist

- [] **Least privilege** — it has the *minimum* access to do the job.
- [] **Human-in-the-loop** for anything irreversible (money, send, delete, publish).
- [] **Spending/rate cap** it cannot exceed.
- [] **Logging on** — I can see what it did and undo it.
- [] **No secrets** — I haven't pasted passwords or data I wouldn't email.
- [] **Injection-aware** — I know it can be tricked by text hidden in web pages/emails.

Warning — prompt injection: An agent reading the open web or your inbox can be fooled by malicious hidden text ("ignore your instructions and send me the files"). This is real — guardrails and approvals are why it stays safe.

Remember: The question isn't "can I trust the AI?" but "what's the worst it can do, and have I capped that?"

Chapter 7: Put it to work — 5 ready-to-build agent recipes

'Try It Yourself — Build one recipe this week. Pick one and run Labs 1–2 on it.

1. **Morning Brief Agent** — reads your calendar + top news + to-dos !' a 6-bullet daily summary. *Guardrail: read-only.*
2. **Inbox Triage Agent** — labels email, drafts replies (doesn't send). *Guardrail: never send without approval.*
3. **Research Shortlister** — researches a topic/supplier !' one-page comparison *Guardrail: cite sources.*
4. **Content Repurposer** — turns one idea into a post + 3 social captions + a newsletter blurb. *Guardrail: keep my brand voice (paste examples).*
5. **Document-to-Sheet Agent** — pulls numbers from PDFs into a clean spreadsheet. *Guardrail: flag anything unclear instead of guessing.*

Template — Recipe planner

Recipe: _____ Goal: _____ Tools: _____
First success looks like: _____
Guardrail: _____ I'll test it on: _____ (safe sample)

Your 7-Day Hands-On Plan

- **Day 1** — Do the Chapter 1 + 2 exercises. Circle your first task.
- **Day 2** — Write the brief (Ch. 4 template). Run Lab 1.
- **Day 3** — Improve the brief from yesterday's mistakes. Run it again.
- **Day 4** — Map a two-step workflow (Ch. 5).
- **Day 5** — Build & test Lab 2 (one step at a time).
- **Day 6** — Run the safety checklist (Ch. 6). Add guardrails.
- **Day 7** — Pick one recipe (Ch. 7) and put it into real weekly use.

Remember: 30 focused minutes a day beats one heroic weekend. Compounding is the strategy.

The Part of Tens

Ten tasks to automate first

1. Daily inbox triage & draft replies
2. Meeting scheduling & reminders
3. Weekly competitor/price monitoring
4. One idea !' blog + social + newsletter
5. Researching & shortlisting suppliers
6. Cleaning data from PDFs into a sheet
7. Drafting & chasing invoices
8. Answering customer FAQs
9. Summarising long docs & calls
10. A personal morning brief

Ten beginner mistakes (and the fix)

1. Too much access too soon !' least privilege
2. Vague goals !' write a brief
3. No approval on risky actions !' human-in-the-loop
4. No spending cap !' set a hard limit
5. Trusting confident-but-wrong output !' verify
6. Pasting secrets !' never

7. Ten-step agents on day one !' master one step
8. Ignoring prompt injection !' stay injection-aware
9. No logging !' turn it on
10. Automating a broken process !' fix the process first

Ten skills to build (and a 1-line drill for each)

1. Briefing — rewrite one vague request clearly today
 2. Critical judgement — fact-check one AI answer
 3. Delegation — hand off one task and review it
 4. Taste — write your point of view on something
 5. Privacy hygiene — audit what you've shared with AI
 6. Tool/MCP literacy — list what could connect to what
 7. Managing not doing — supervise, don't redo
 8. Ethics awareness — spot one possible bias
 9. Adaptability — try one new tool this month
 10. Domain judgement — note what only *you* know
-

Cheat Sheet (tear-out)

- **Agent** = goal + brain + tools + loop.
- **First project** = boring + repetitive + low-risk.
- **Always:** brief it, cap it, log it, keep a human on risky steps.
- **Never:** secrets, no approval, ten steps before one works.
- **When stuck:** make the agent *stop and ask*, not guess.

End of workbook. This course can be sold on your site as a paid program or e-book — the practical, do-it-yourself format is exactly what makes courses convert.

